RAG Pipeline Overview

1. Load documents (PDFs, text files, etc.)

2. Chunk documents

3. Generate embeddings

4. Retrieve relevant chunks for a query